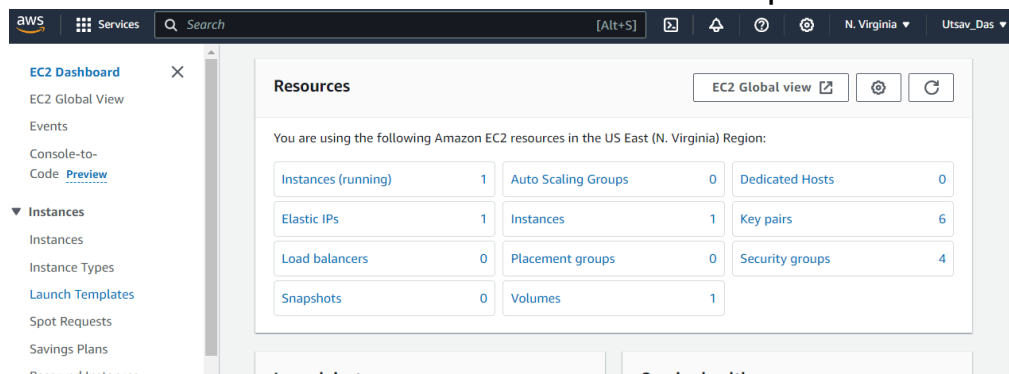


Assignment – 11

Problem Statement:

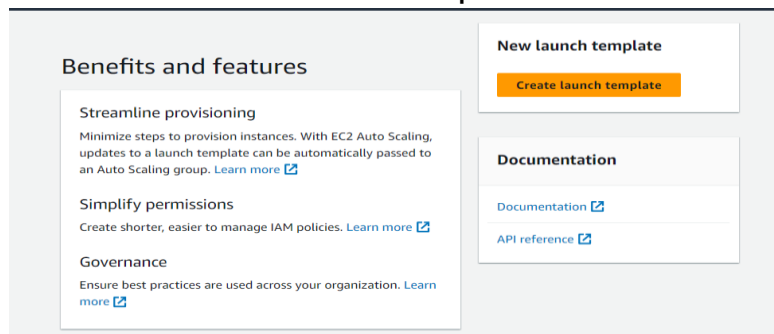
Build scaling plans in AWS that balance the load on different EC2 instances.

1) Go to EC2 then in Instances click on Launch Templates.

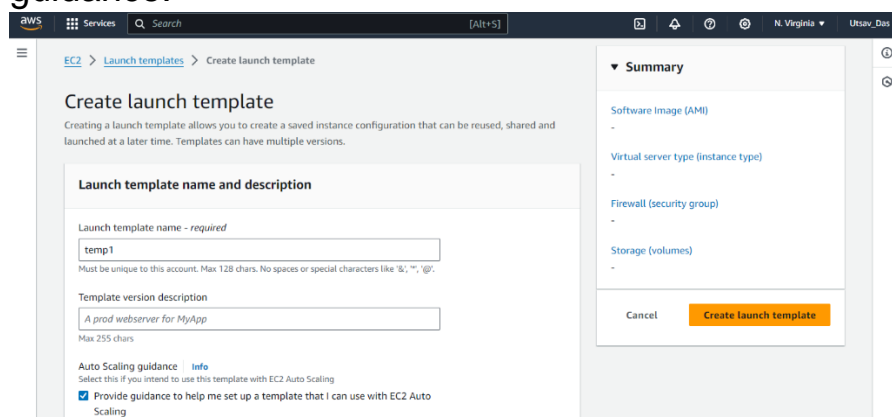


48

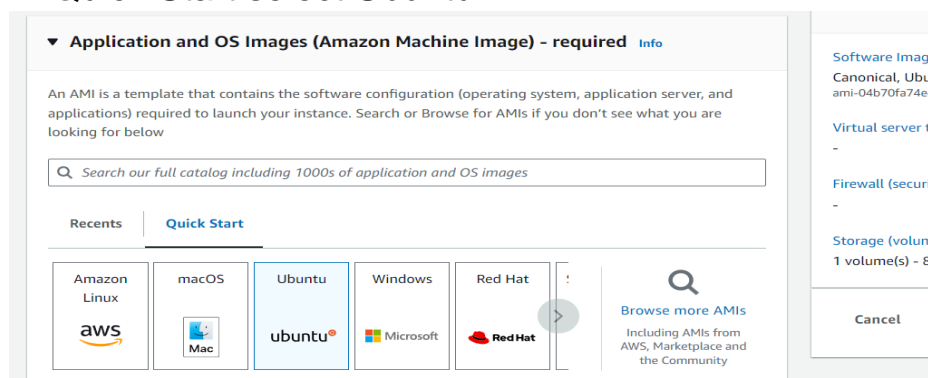
2) Click on Create launch template.



3) Give Launch template name and then click on checkbox for Auto Scaling guidance.



4) In Quick Start select Ubuntu.



5) Select Instance type – either t2.micro or t3.micro. Create new key pair.

The screenshot shows the AWS EC2 console configuration page. Under the 'Instance type' section, 't2.micro' is selected, which is marked as 'Free tier eligible'. Details for t2.micro include: Family: t2, 1 vCPU, 1 GiB Memory, Current generation: true, On-Demand Windows base pricing: 0.0162 USD per Hour, On-Demand SUSE base pricing: 0.0116 USD per Hour, On-Demand RHEL base pricing: 0.0716 USD per Hour, and On-Demand Linux base pricing: 0.0116 USD per Hour. There is a toggle for 'All generations' and a link to 'Compare instance types'. Below this, a note states 'Additional costs apply for AMIs with pre-installed software'. The 'Key pair (login)' section explains that a key pair is used to securely connect to the instance. A dropdown menu for 'Key pair name' shows 'key-new2', and there is a 'Create new key pair' button.

6) Select existing security group.

The screenshot shows the 'Network settings' section of the AWS EC2 console. The 'Subnet' dropdown is set to 'Don't include in launch template', with a link to 'Create new subnet'. A note states: 'When you specify a subnet, a network interface is automatically added to your template.' The 'Firewall (security groups)' section has an 'Info' link and explains that a security group is a set of firewall rules. There are two radio buttons: 'Select existing security group' (which is selected) and 'Create security group'. Below this, the 'Security groups' dropdown shows 'Select security groups'. A list of security groups is displayed, including 'my-security-group-1' with ID 'sg-09c311ab76a75378b' and VPC 'vpc-05e13defc1bb34724'. There is a 'Compare security group rules' link. At the bottom, there is a link for 'Advanced network configuration'.

7) Now go to Advanced details and then in User data write given statements and then Create launch template.

```
#!/bin/bash
apt -get update
apt -get install -y nginx
systemctl start nginx
systemctl enable nginx
apt -get install -y git
curl -SL https://deb.nodesource.com/setup_16.x|sudo -E bash -
apt -get install -y nodejs
git clone https://github.com/UnderDevelopment10/new-repo1.git
cd new-repo1
npm install
node index.js
```

Allow tags in metadata [Info](#)

Don't include in launch template

User data - optional [Info](#)

Upload a file with your user data or enter it in the field.

[Choose file](#)

```
#!/bin/bash
apt-get update
apt-get install -y nginx
systemctl start nginx
systemctl enable nginx
apt-get install -y git
curl -SL https://deb.nodesource.com/setup_16.x | sudo -E bash -
apt-get install -y nodejs
git clone https://github.com/UnderDevelopment10/new-repo1.git
cd new-repo1
npm install
node index.js
```

☐ User data has already been base64 encoded

Summary

Software Image (AMI)
Canonical, Ubuntu, 24.04 LTS, ...[read more](#)
ami-04b70fa74e45c3917

Virtual server type (instance type)
t2.micro

Firewall (security group)
my-security-group-1

Storage (volumes)
1 volume(s) - 8 GiB

Cancel [Create launch template](#)

- 8) After creating Launch template, click on Auto Scaling Groups in left pane. Click on Create Auto Scaling group.

AMIs

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Load Balancing

Load Balancers

Target Groups

Trust Stores [New](#)

Auto Scaling

[Auto Scaling Groups](#)

Amazon EC2 Auto Scaling

helps maintain the availability of your applications

Auto Scaling groups are collections of Amazon EC2 instances that enable automatic scaling and fleet management features. These features help you maintain the health and availability of your applications.

[How it works](#)

[Pricing](#)

Amazon EC2 Auto Scaling features have no additional fees beyond the service fees

Create Auto Scaling group

Get started with EC2 Auto Scaling by creating an Auto Scaling group.

[Create Auto Scaling group](#)

- 9) Give Auto scaling group name and then select your recently launched template. Then click on Next.

Choose launch template [Info](#)

Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group.

Name

Auto Scaling group name

Enter a name to identify the group.

myautoscale1

Must be unique to this account in the current Region and no more than 255 characters.

Launch template [Info](#)

For accounts created after May 31, 2023, the EC2 console only supports creating Auto Scaling groups with launch templates. Creating Auto Scaling groups with launch configurations is not recommended but still available via the CLI and API until December 31, 2023.

Launch template

Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.

temp1

[Create a launch template](#)

- 10) Select availability zones and subnets then click on Next.

Availability Zones and subnets

Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets

us-east-1a | subnet-065f62c266b1adeae X

172.31.80.0/20 Default

us-east-1b | subnet-05a5c4402a0ca882f X

172.31.16.0/20 Default

us-east-1c | subnet-0bc2b74d3abb99de1 X

172.31.32.0/20 Default

us-east-1d | subnet-048e86a467c30c0bd X

172.31.0.0/20 Default

us-east-1e | subnet-034e9f95c5897aac3 X

172.31.48.0/20 Default

us-east-1f | subnet-04093473c47d70dab X

172.31.64.0/20 Default

[Create a subnet](#)

Cancel [Skip to review](#) [Previous](#) [Next](#)

11) Click on Attach to a new load balancer.

Configure advanced options - optional [Info](#)

Integrate your Auto Scaling group with other services to distribute network traffic across multiple servers using a load balancer or to establish service-to-service communications using VPC Lattice. You can also set options that give you more control over health check replacements and monitoring.

Load balancing [Info](#)

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

☐ **No load balancer**
Traffic to your Auto Scaling group will not be fronted by a load balancer.

☐ **Attach to an existing load balancer**
Choose from your existing load balancers.

☒ **Attach to a new load balancer**
Quickly create a basic load balancer to attach to your Auto Scaling group.

12) Select Application Load Balancer as load balancer type and Internet-facing as Load balancer scheme.

Attach to a new load balancer
Define a new load balancer to create for attachment to this Auto Scaling group.

Load balancer type
Choose from the load balancer types offered below. Type selection cannot be changed after the load balancer is created. If you need a different type of load balancer than those offered here, [visit the Load Balancing console](#).

☒ **Application Load Balancer**
HTTP, HTTPS

☐ **Network Load Balancer**
TCP, UDP, TLS

Load balancer name
Name cannot be changed after the load balancer is created.

Load balancer scheme
Scheme cannot be changed after the load balancer is created.

☐ Internal

☒ Internet-facing

Network mapping

13) Select port no. 4000 for HTTP and select New target group name.

☒ us-east-1f

☒ us-east-1a

Listeners and routing
If you require secure listeners, or multiple listeners, you can configure them from the [Load Balancing console](#) after your load balancer is created.

Protocol	Port	Default routing (forward to)
HTTP	<input type="text" value="4000"/>	<input type="text" value="Create a target group"/>

New target group name
An instance target group with default settings will be created.

Tags - optional
Consider adding tags to your load balancer. Tags enable you to categorize your AWS resources so you can more easily manage them.

50 remaining

14) Now turn on Elastic load balancing health check and in health check grace period give 240 seconds. Then click on Next.

Health checks
Health checks increase availability by replacing unhealthy instances. When you use multiple health checks, all are evaluated, and if at least one fails, instance replacement occurs.

EC2 health checks
[Always enabled](#)

Additional health check types - optional [Info](#)

☒ **Turn on Elastic Load Balancing health checks** **Recommended**
Elastic Load Balancing monitors whether instances are available to handle requests. When it reports an unhealthy instance, EC2 Auto Scaling can replace it on its next periodic check.

☒ **EC2 Auto Scaling will start to detect and act on health checks performed by Elastic Load Balancing.**
To avoid unexpected terminations, first verify the settings of these health checks in the [Load Balancer console](#).

☐ **Turn on VPC Lattice health checks**
VPC Lattice can monitor whether instances are available to handle requests. If it considers a target as failed a health check, EC2 Auto Scaling replaces it after its next periodic check.

Health check grace period [Info](#)
This time period delays the first health check until your instances finish initializing. It doesn't prevent an instance from terminating when placed into a non-running state.
 seconds

15) Select 2 in Desired capacity in group size and in Scaling 2 as Min desired capacity and 3 as Max desired capacity.

Group size [Info](#)
Set the initial size of the Auto Scaling group. After creating the group, you can change its size to meet demand, either manually or by using automatic scaling.

Desired capacity type
Choose the unit of measurement for the desired capacity value. vCPUs and Memory(GiB) are only supported for mixed instances groups configured with a set of instance attributes.

Units (number of instances) ▼

Desired capacity
Specify your group size.

2

Scaling [Info](#)
You can resize your Auto Scaling group manually or automatically to meet changes in demand.

Scaling limits
Set limits on how much your desired capacity can be increased or decreased.

Min desired capacity **Max desired capacity**

2 3

Equal or less than desired capacity Equal or greater than desired capacity

16) Next in Automatic scaling select Target tracking scaling policy and give 240 in Instance warmup.

Automatic scaling - optional [Info](#)
Choose whether to use a target tracking policy. You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.

☐ No scaling policies
Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.

☒ **Target tracking scaling policy**
Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.

Scaling policy name

Target Tracking Policy

Metric type [Info](#)
Monitored metric that determines if resource utilization is too low or high. If using EC2 metrics, consider enabling detailed monitoring for better scaling performance.

Average CPU utilization ▼

Target value

50

Instance warmup [Info](#)

240 seconds

☐ Disable scale in to create only a scale-out policy

17) Then click on Next, then again next.

Instance scale-in protection
Scale-in protection prevents newly launched instances from being terminated by scaling activities. Make sure to remove scale-in protection for the group or individual instances when instances are ready to be terminated.

☐ Enable instance scale-in protection

Add notifications - optional [Info](#)
Send notifications to SNS topics whenever Amazon EC2 Auto Scaling launches or terminates the EC2 instances in your Auto Scaling group.

Add notification

Add tags - optional [Info](#)
Add tags to help you search, filter, and track your Auto Scaling group across AWS. You can also choose to automatically add these tags to instances when they are launched.

☒ You can optionally choose to add tags to instances (and their attached EBS volumes) by specifying tags in your launch template. We recommend caution, however, because the tag values for instances from your launch template will be overridden if there are any duplicate keys specified for the Auto Scaling group.

Tags (0)

Add tag

10 remaining

18) Now click on Create Auto Scaling group.

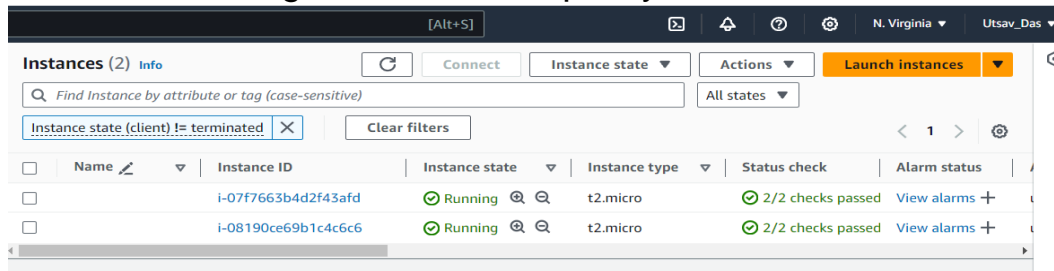
Step 6: Add tags [Edit](#)

Tags (0)

Key	Value	Tag new instances
No tags		

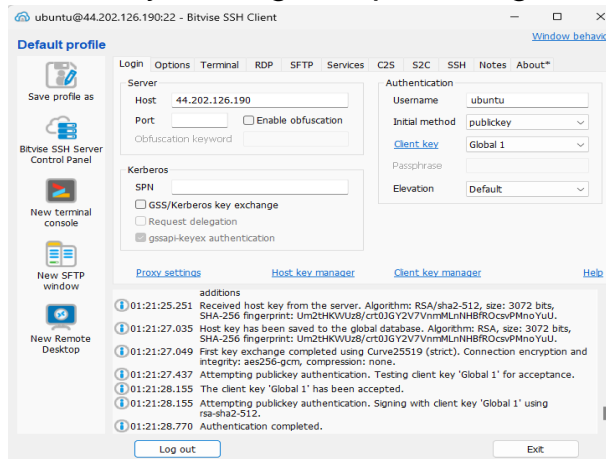
Cancel Previous **Create Auto Scaling group**

- 19) Come back to EC2 dashboard and go to Instance and here you can see 2 instances running as minimum capacity 2 is chosen.




- 20) Click on any instance and then copy public IPv4 address.

- 21) Open Bitwise SSH client and then paste that public IPv4 address and then in client key manager import the generated key and then Login.



- 22) Open New terminal console and then write command 'sudo nano infy.sh'. A new .sh file will be created. Now write this code to run an infinite loop.

```
#!/bin/bash
while(true)
do
    echo "Inside Loop"
done
```




- 23) Then do Ctrl+X then Y then click Enter. And write command 'sudo chmod 777 infy.sh' to provide all permissions to file and then to run give 'sh infy.sh'.

```
ubuntu@ip-172-31-83-21:~$ sudo nano infy.sh
ubuntu@ip-172-31-83-21:~$ sudo chmod 777 infy.sh
ubuntu@ip-172-31-83-21:~$ sh infy.sh
```

- 24) Here we can see infinite loop running.

```
Inside Loop
Inside Loop
Inside Loop
Inside Loop
Inside Loop
```

25) Go back to instances and then select both instances and then in bottom pane, go to CPU utilization and select Enlarge.



26) The graph displays CPU utilization for both instances. When CPU utilization for one instance is very high, then another instance will be created as we have set maximum capacity to 3.

27) Finally in Instances, we can see another new instance is created.

